

Integrating OpenTelemetry & Security in eShop

1. Objective

1. **Implement OpenTelemetry tracing** on a single feature or use-case (end-to-end).
2. **Mask or exclude sensitive data** (e.g., email, payment details) from telemetry and logs.
3. **Set up a basic Grafana dashboard** to visualize the traces and metrics.
4. (Optional Extras) **Explore data encryption and compliance** in the database layer and introduce **column masking** for sensitive data.

2. Project Overview

eShop is a modern e-commerce system built using ASP.NET Core and microservices. The official repo can be found here:

- [dotnet/eShop GitHub Repository](https://github.com/dotnet/eShop)

Key aspects of eShop include:

- **Microservices** split by bounded contexts (e.g., Ordering, Catalog, Basket).
- **Event-driven communication** using message brokers.
- **Containerization** with Docker.
- Multiple **frontend options** (Blazor, MVC, Angular, MAUI).
- Uses **IdentityServer** for authentication (though this project scope doesn't extend into deeper Auth improvements).

Students will **fork** this repository and focus on adding **OpenTelemetry** instrumentation for just **one** selected feature or user flow, rather than instrumenting the entire system.

3. Student Tasks

3.1 Integrate OpenTelemetry Tracing (Single Feature)

1. **Pick a Feature**

- Identify **one** user flow, such as “Place an Order” or “Add to Cart.”
- Ensure all calls (HTTP, database, messaging) within this feature are traced end-to-end.

2. Instrument Services

- Add OpenTelemetry libraries (e.g., [OpenTelemetry.Instrumentation.AspNetCore](#) and [OpenTelemetry.Instrumentation.Http](#)).
- Wrap key service calls with **Activity** or built-in instrumentation to collect spans.

3. Data Scrubbing

- **Mask or exclude sensitive data** like user emails and payment details from traces and logs.
- Validate your logs and traces to ensure no personal data is leaked.

4. Exporters

- Configure an exporter (e.g., **Jaeger** or **Prometheus**).
- This feeds your traces/metrics into tools that you can visualize in **Grafana**.

3.2 Observability Setup

1. Metrics

- Expose basic metrics (e.g., request duration, success/error rates).
- Use an OpenTelemetry collector or a built-in ASP.NET Core metrics endpoint.

2. Grafana Dashboard

- Install or configure **Grafana** to visualize traces or metrics.
- Demonstrate how to filter and analyze the single feature flow you instrumented.

3. Simulate Load

- Create a simple load test (using **curl**, **k6**, **Locust**, or **JMeter**) that triggers the selected feature repeatedly.
- Observe how traces and metrics appear in real time on Grafana.

3.3 (Optional) Data Security & Compliance

1. Column Masking Challenge

- **Partially or fully mask** columns containing sensitive data (e.g., credit cards, emails).
- Use either SQL Server Dynamic Data Masking (if using SQL Server) or custom masking logic in your data-access layer.
- Make sure your masked fields do not appear in logs/traces in plaintext.

4. Deliverables

1. Source Code

- Forked eShop repository with instrumentation code for the single feature.
- Updated configuration for OpenTelemetry (tracing + metrics) and any code for masking sensitive data.
- (Optional) Data encryption or column masking logic.

2. Documentation

- **README.md** with:
 - How to build and run the eShop environment with your changes.
 - Steps to configure and launch the OpenTelemetry collectors/exporters.
 - Instructions to set up or view the Grafana dashboard.
- Basic **Architecture Diagram** or sequence diagram highlighting the traced feature.

3. Grafana Dashboard

- A dashboard showing traces/metrics for the selected feature.
- Screenshots or instructions for easy verification.

4. Load Test Scripts

- Provide or link to a script (k6, JMeter, etc.) that generates activity on your selected feature.
- This demonstrates how data flows end-to-end while being traced.

5. Evaluation Criteria

1. Correctness & Completeness

- The chosen feature is properly traced end-to-end (frontend → microservice → database/message broker).

- Sensitive data (emails, payment info) is masked in logs/traces.

2. Observability Setup

- Grafana dashboard effectively visualizes request flows and key metrics.
- Students can demonstrate how to debug or analyze performance using these traces/metrics.

3. Security & Compliance (Optional)

- If attempted: column masking or data encryption is implemented according to best practices.
- No sensitive data leaks in logs, traces, or in the database output for non-privileged users.

4. Documentation & Clarity

- README and architecture notes are clear and concise.
- The single feature flow is explained, including why it was chosen and how it was instrumented.

6. Hints & Tips

- OpenTelemetry Setup

- Start with a minimal approach: instrument HTTP client and server using built-in OTel packages.
- Use [ActivitySource](#) for custom spans if you need to trace specific logic inside your code.

- Data Scrubbing

- Check logs (Serilog, NLog, or [.NET ILogger](#)) for any personal data.
- Implement your own log filter or OTel processor to remove PII from spans.

- Grafana Dashboards

- If using Jaeger, you can embed Jaeger panels in Grafana, or use the Prometheus data source for metrics.
- Make your panels user-friendly and label them to match your selected feature.

- Load Testing

- Even a simple shell script that sends repeated HTTP requests can give you enough volume for demonstration.
 - Combine load generation with a break-it test approach (stop a service, observe logs/traces, etc.).
- **Column Masking**
- If you choose to do this, experiment with partial masking: e.g., show only the first 3 letters of email.
 - Make sure you have a clear policy on who can see the real data (if implementing role-based unmasking).
 - SQLServer supports [Dynamic Data Masking](#).
 - Postgres: <https://www.bytebase.com/blog/postgres-data-masking/> or https://postgresql-anonymizer.readthedocs.io/en/latest/dynamic_masking/
 -
 - Question: can the DB Admins see the unmasked information?